

Homework 5

All assignments need to be submitted via github classroom:

and on gradescope. You can submit in groups of maximum size two.

The homework is due 05/01/19 at 1pm.

All the tasks are to be completed using the [keras Sequential interface](#). It's recommended that you run your code on GPU using google colab:

<https://colab.research.google.com/>

You can enable GPU support at “runtime” -> “change runtime type”.

Feel free to experiment with TPU support if you're adventurous.

You can use any other resources at your disposal if you prefer. You should not use k-fold cross-validation for any of these tasks. You can use StratifiedShuffleSplit to create a single train-validation split for use with GridSearchCV.

Use of GridSearchCV might not be the best option for any task but task1, though.

Task 1 [10 Points]

Run a multilayer perceptron (feed forward neural network) with two hidden layers and rectified linear nonlinearities on the iris dataset using the keras [Sequential interface](#). Include code for selecting regularization strength and number of hidden units using GridSearchCV and evaluation on an independent test-set.

Task 2 [40 Points]

Train a multilayer perceptron on the Fashion MNIST dataset using the traditional train/test split as given by `fashion_mnist.load_data` in keras. Use a separate 10000 samples (from the training set) for model selection and to compute learning curves (accuracy vs epochs, not vs `n_samples`). Compare a “vanilla” model with a model using drop-out (potentially a bigger model), and to a model using batch normalization and residual connections (but not dropout). Visualize learning curves for all models.

Task 3 [60 Points]

Train a convolutional neural network on the following dataset:

<https://www.kaggle.com/paultimothymooney/breast-histopathology-images>

- 3.1** Start with a model without residual connections (using batch normalization is likely to be helpful and you should try it, whether you use dropout is your choice).
- 3.2** Augment the data using rotations, mirroring and possibly other transformations. How much can you improve your original model by data augmentation?
- 3.3** Build a deeper model using residual connections. Show that you can build a deep model that would not be able to learn if you remove the residual connections (i.e. compare a deep model with and without residual connections while the rest of the architecture is constant).

Hint: Make sure you are doing the reshape for the training set correctly. A direct reshape might give you garbled images. Display an image after reshaping to make sure they are correct.

Some additional advice to help you along:

- Make sure all your code is running on GPU. Use `"sess = tf.Session(config=tf.ConfigProto(log_device_placement=True))"` to start the tensorflow session to see which device is being used and confirm it is the device you intended.
- Preprocess the images before training a model.
- Test your code on a small part of the data before training the model. You don't want your code to fail on a print statement after waiting for the network to train.